

# Virtualization and Containers

Open Source Software

Wesley Turner

Rensselaer Polytechnic Institute

Department of Computer Science

Licensed under: CC-BY-SA 4.0 <https://creativecommons.org/licenses/by-sa/4.0/>

## Reading Material

### Reading Material

- **Virtualization** – <https://en.wikipedia.org/wiki/Virtualization>  
– [https://en.wikipedia.org/wiki/Hardware\\_virtualization](https://en.wikipedia.org/wiki/Hardware_virtualization)
- **Using Virtual Machines and Docker:** – <https://goo.gl/sRVT7o>
- **Oracle VirtualBox VM (Reference)** – <https://www.virtualbox.org/manual/ch01.html#virt-why-useful>

## Introduction

### What is Virtualization?

We will consider two main types:

- Virtual Machines
  - Software to allow a piece of hardware to run multiple operating system images at the same time
  - Eg. *VirtualBox*
- Containers
  - A lightweight, stand-alone, executable package of a piece of software that includes everything needed to run it
  - Eg. *Docker*

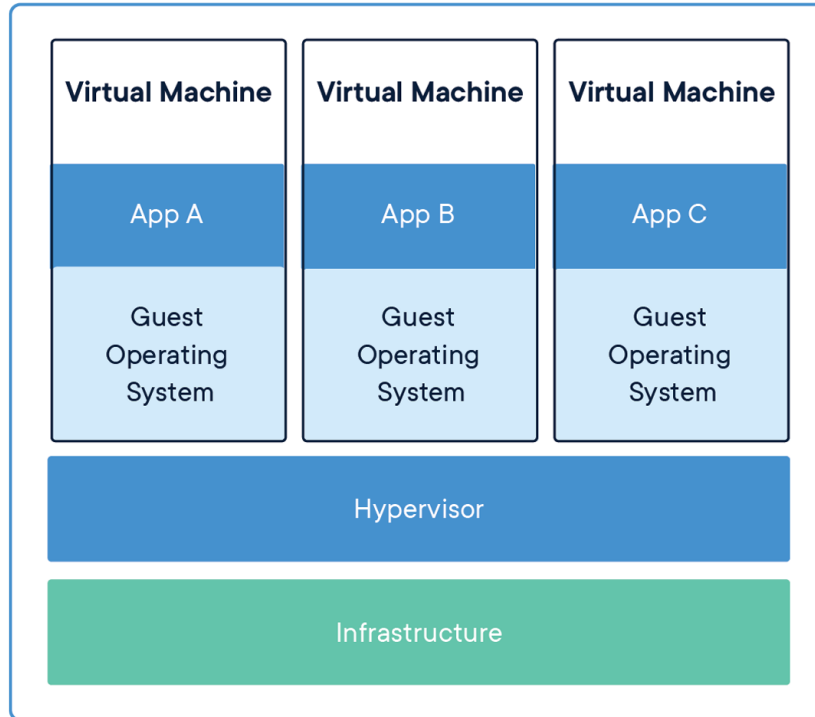
## Virtual Machines

### Concept

In a Virtual Machine:

- A software layer (*hypervisor*) emulates the underlying hardware
- An additional operating system runs on the virtual hardware
- Users interact with the virtualization layer as if it were the underlying machine

## Architecture



[https://www.docker.com/resources/what-container/#/virtual\\_machines](https://www.docker.com/resources/what-container/#/virtual_machines)

## Advantages

- Virtual Machines share resources
  - Memory, disk, network bandwidth
  - Better resource utilization
- They run in isolated environments
  - Better security
- The same hardware can run multiple operating systems
  - Eg. Windows and Linux can coexist on the same machine at the same time
- An entire environment can be distributed
  - Less time wasted setting up environments
  - Less confusion about versions or compatibilities
- Machine states can be checkpointed
  - Can be migrated to a more or less powerful machine as needed
  - Rapid return of services when hardware goes down

## Disadvantages

- Virtual Machines share resources
  - Memory, disk, network bandwidth
  - Poor planning and multiple VMs can drown a machine
- The same hardware can run multiple operating systems
  - May require multiple licenses
  - A problem with one set of hardware can affect hundreds of VMs
- An entire environment can be distributed
  - The entire guest environment -- including the entire guest OS, disk and memory -- needs to be represented in the VM file
    - \* VMs are large
  - Each environment communicates to resources through the hypervisor
    - \* Guest OS <-> hypervisor <-> native OS
      - Can significantly impact performance
- Machine states can be checkpointed
  - Checkpoint files can be huge

## Virtualization and the Cloud

Virtual machine are the *de facto* way of using cloud compute resources

- **You** are protected by the virtual environment by interference from other VMs
- Other VMs, and the provider, are protected from **you** as well
- You can easily add VMs or migrate to more/less powerful VMs as your needs change

You do not necessarily know what hardware you are running on

- Performance can be different on different systems
- Performance can be different on different clusters
- Performance can be different in different locations (proximity sometimes matters)

## An Example

### The VA Scheduling Competition

An interesting cloud use for Virtual Machines: <https://vascheduling.devpost.com/>

- The VA was investigating a new scheduler for their Electronic Health Recorder (EHR) system
  - Contestants were each given a Virtual Machine with access to nominal (simulated) patient data and a version of a running VISTA EHR
    - \* The machines ran in the cloud on Rackspace <https://www.rackspace.com/>

- The VA scheduler (continued):
  - They developed their scheduling solution on the system
  - At the end of the competition, the machines were checkpointed and submitted to OSHERA (Open Source Electronic Health Record Agent) for evaluation
- Three awards were given:
  1. <https://devpost.com/software/health-etime>
  2. <https://devpost.com/software/oh-scheduler>
  3. <https://devpost.com/software/hp-open-community-team-submission>

## VirtualBox

### VirtualBox



<https://www.virtualbox.org/>

Virtualbox is supported by Oracle and comes with a full suite of documentation and other useful information

- Licensing:
  - GNU General Public License (GPL) version 2
- Documents:
  - <https://www.virtualbox.org/manual/UserManual.html>
- Getting VirtualBox:
  - <https://www.virtualbox.org/wiki/Downloads>
- Source

`svn co https://www.virtualbox.org/svn/vbox/trunk vbox`

- Contributing
  - [https://www.virtualbox.org/wiki/Contributor\\_information](https://www.virtualbox.org/wiki/Contributor_information)
- Community
  - <https://www.virtualbox.org/wiki/Community>
- Issue Tracker
  - <https://www.virtualbox.org/ticket/20939>
- Runs on:
  - Windows, Linux, Macintosh, and Solaris
- Guest OS:
  - Windows (NT 4.0, 2000, XP, Server 2003, Vista, Windows 7, Windows 8, Windows 10)
  - DOS/Windows 3.x
  - Linux (2.4, 2.6, 3.x and 4.x)
  - Solaris
- Guest OS continued:
  - OpenSolaris
  - OS/2
  - OpenBSD
  - Other ... [https://www.virtualbox.org/wiki/Guest\\_OSes](https://www.virtualbox.org/wiki/Guest_OSes)
- Licensing:
  - VirtualBox is Open Source, but you still need permission for any guest operating system you use
  - In particular, Windows requires appropriate licensing

A simple interface can be used to define your virtual machine



- The image can then be launched using the Virtualbox program
  - It gives a window where the entire guest operating system can be accessed
- Other buttons allow you to stop and checkpoint (save) the current state
- Be careful provisioning your image:
  - There is an inherent trade-off between image capabilities (size/memory) and VM compactness
  - That said, changing the size can be a pain
  - In general, you want to use the smallest VM that meets your needs

## Containers

### Concept

In a Container, the abstraction is to the application layer instead of the hardware layer:

- The application layer interacts directly with a *Container* layer
  - There is no Guest OS layer
  - The guest layer capabilities come in as shared services in the docker layer
  - The containers share the operating system kernel
- Each container runs as an application in user space

- Containers are stateless and use the existing system resources to store data and state

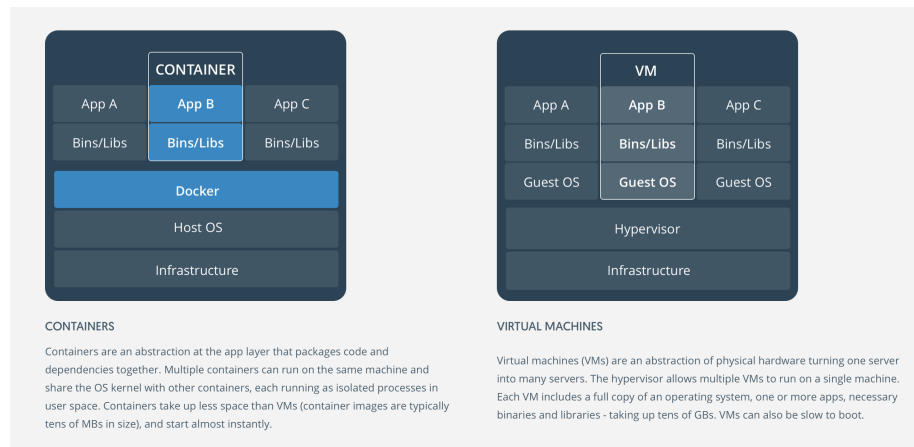
## Advantages

- Containers were designed to solve some of the issues of VMs. Paticularly,
  - Large image sizes
  - High latency
  - Reduced performance
- By linking directly to the host kernel, containers:
  - Can have a reduced size
  - Can share resources
  - Can have improved performance

## Disadvantages

Since containers are built from the bottom up, they can be very compact and light weight, but that means that you need to install/configure everything you want to use in your container.

## Architecture



[https://www.docker.com/what-container#/virtual\\_machines](https://www.docker.com/what-container#/virtual_machines)

<https://docs.microsoft.com/en-us/virtualization/windowscontainers/about/containers-vs-vm>

## Docker

### Docker Images

Docker maintains a library of images that you can run, i.e.:

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
python	3.5	99034fcd3ca6	2 weeks ago	917MB
python	latest	746a826ed9d7	2 weeks ago	922MB
redis	latest	f06a5773f01e	2 weeks ago	83.4MB
node	8	ed145ef978c4	2 weeks ago	673MB
postgres	9.5	a4da5bcea663	2 weeks ago	234MB
mongo	3.2	386ac407ef9f	5 weeks ago	300MB
hello-world	latest	e38bc07ac18e	3 months ago	1.85kB
ubuntu	16.04	20c44cd7596f	8 months ago	123MB
ubuntu	14.04	d6ed29ffda6b	8 months ago	221MB
fedora	latest	422dc563ca32	8 months ago	252MB
mysql	latest	5709795eeffa	9 months ago	408MB
ubuntu	latest	dd6f76d9cc90	9 months ago	122MB
node	7-alpine	4b72b56791f9	13 months ago	58.3MB
codenvy/ubuntu_jdk8	latest	4074bfc5705b	20 months ago	668MB
docker/whalesay	latest	6b362a9f73eb	3 years ago	247MB

```
$ docker run -it docker/whalesay
root@40aa8aad15f:/cowsay# cowsay squirt
```

```

-----
< squirt >
-----

```

You can find out what images you have running with

```
Wesleys-MacBook-Pro:Virtualization_Lecture wes$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
eec278129a30	mongodev	"/bin/bash"	5 days ago	Up 5 days		zen williams

1. Create a DockerFile to define your image
2. Run a docker build to turn your configuration into an image
3. Run just like a pre-canned image



4. Once you have build a docker image, it will appear when you run *docker images*
5. Alternately, you can use a yaml file to define multiple images that work together to create a system

We will explore all of these in the following example.

## Example

### Quick Cheat Sheet

```
## List Docker CLI commands
docker
docker container --help

## Display Docker version and info
docker --version
docker version
docker info

## Execute Docker image
docker run hello-world

## List Docker images
docker image ls

## List Docker containers (running, all, all in quiet mode)
docker container ls
docker container ls --all
docker container ls -aq
```

## Example

As a class exercise, we will run through steps 1-4 of the Docker Tutorial at: <https://docs.docker.com/get-started/>

- Errata
  - Where the Dockerfile says 12-alpine, use 14-alpine
  - Where the Dockerfile says python2, use python3

For those of you running on WSL, there may be a few more steps:

- Follow the directions to download the Windows Desktop
- Then go to: <https://docs.microsoft.com/en-us/windows/wsl/install-win10>
- And Finally: <https://docs.docker.com/docker-for-windows/wsl/>

## End